

# SPARC Containerization and Deployment


## 1.0 Introduction

This notebook covers the final phase: packaging the SPARC backend into portable containers and deploying them to HiPerGator with a robust networking bridge.

### 1.1 Objectives

1. **Containerize**: Create Dockerfiles for the Multi-Agent System (MAS).
2. **Bridge**: Configure the WebSocket-to-gRPC bridge for Unity connectivity.
3. **Deploy**: Generate production SLURM scripts for HiPerGator.

### 1.2 Introduction Diagram




Introduction: This section sets the objectives for packaging and deploying the backend. The goal is to containerize the Multi-Agent System (MAS), configure the WebSocket-to-gRPC bridge for Unity connectivity, and generate robust production SLURM scripts for deployment to HiPerGator.

## 2.0 Containerization (Docker -> Apptainer)

We develop with Docker/Podman and deploy with Apptainer on HPC.

### 2.0 Container Build Strategy Diagram



**Container Build Strategy:** This flow shows the Multi-Stage Build strategy used to create secure and small containers. A "Builder" stage uses Poetry to compile dependencies, and then only the necessary artifacts are copied over to a slim "Runtime" stage. This excludes compiler tools and cache files from the final production image.

## 2.1 Dockerfile Definition

This script creates a `Dockerfile.mas` for the Multi-Agent System. We use a multi-stage build strategy:

1. **Builder Stage:** Installs dependencies using Poetry.
2. **Runtime Stage:** Copies only the installed packages to a lightweight `python:3.10-slim` image. This minimizes the container size and attack surface.

## 2.2 Dockerfile for Multi-Agent System (MAS)

```

# 2.1 Dockerfile for Multi-Agent System (MAS)

def create_dockerfile():
    dockerfile_content = """
# --- Build Stage ---
FROM python:3.10-slim as builder
WORKDIR /app
RUN pip install --no-cache-dir poetry
COPY poetry.lock pyproject.toml ./
RUN poetry config virtualenvs.create false && poetry install --no-dev --no-interaction --no-ansi
COPY . .

# --- Runtime Stage ---
FROM python:3.10-slim
WORKDIR /app
COPY --from=builder /usr/local/lib/python3.10/site-packages
/usr/local/lib/python3.10/site-packages
COPY --from=builder /app /app
EXPOSE 8000
CMD ["uvicorn", "main:app", "--host", "0.0.0.0", "--port", "8000"]
"""


    with open("Dockerfile.mas", "w") as f:
        f.write(dockerfile_content.strip())
    print("Created Dockerfile.mas")
  
```

```
create_dockerfile()
```

## 3.0 Local Development with Podman

Podman allows creating a 'pod' to simulate the production network namespace.

### 3.0 Local Development Pod Diagram



**Local Development Pod (Podman):** This illustrates the local development environment using Podman Pods. Unlike standard Docker containers which are isolated, a "Pod" shares a network namespace (localhost). This allows the Riva Server, WebSocket Bridge, and MAS (Multi-Agent System) to communicate locally, perfectly simulating the production environment on a developer's machine.

### 3.1 Podman Local Workflow

For local development, **Podman** is preferred over Docker because it allows us to create a **Pod**. A Pod shares a network namespace (localhost), allowing the separate containers (Riva, Bridge, MAS) to communicate with each other as if they were running on the same machine, mimicking the production environment.

### 3.2 Podman Workflow (Reference Commands)

```
# 3.1 Podman Workflow (Reference Commands)
# Run these in your local terminal to test interaction between Riva, Bridge, and
MAS.

podman_commands = """
# 1. Create Pod
podman pod create --name sparc-backend -p 8080:8080

# 2. Run Riva Server
podman run -d --pod sparc-backend --name riva-server nvcr.io/nvidia/riva/riva-
```

speech:2.16.0-server

```
# 3. Run WebSocket Bridge
podman run -d --pod sparc-backend --name ws-bridge \
-e RIVA_API_URL=localhost:50051 \
riva-websocket-bridge:latest


# 4. Run MAS Server
podman run -d --pod sparc-backend --name mas-server your-repo/mas-server:latest
"""

print(podman_commands)
```

## 4.0 Production Deployment on HiPerGator

Deploying persistent services using SLURM and Apptainer.

### 4.0 Production Deployment Diagram



**Production Deployment (SLURM):** This diagram shows the execution flow of the `sparc_production.slurm` script on HiPerGator. It details how the SLURM scheduler allocates resources (GPUs) and then launches three concurrent Apptainer containers in the background, keeping them alive with a wait command.

### 4.1 Building SIF Images

HiPerGator uses Apptainer, which requires Singularity Image Format (`.sif`) files. The commands below (commented out) show how to convert your local Docker images into SIF files using `apptainer build`. These files should be stored in the `/blue` directory.

## 4.2 Build SIF Images

```
# 4.1 Build SIF Images
# !module load apptainer
# !apptainer build mas_server.sif docker-daemon://your-repo/mas-server:latest
# !apptainer build websocket_bridge.sif docker-daemon://riva-websocket-
bridge:latest
print("Build SIF images from Docker/Daemon sources before deployment.")
```

## 4.3 Production Service Launch

This function generates the `sparc_production.slurm` script. This is the critical deployment artifact that runs the system on HiPerGator. Key features:

- **Persistent GPUs:** Requests 4 GPUs on the AI partition.
- **Background Processes:** Launches Riva, the Bridge, and the MAS server as background tasks (`&`).
- **Wait Command:** The `wait` instruction keeps the SLURM job alive indefinitely, ensuring the services remain running.

## 4.4 Production SLURM Script Generator

```
# 4.2 Production SLURM Script Generator

def generate_production_script():
    script_content = """
#!/bin/bash
#SBATCH --job-name=sparc-production-service
#SBATCH --partition=hpg-ai
#SBATCH --nodes=1
#SBATCH --gpus-per-task=4
#SBATCH --cpus-per-task=32
#SBATCH --mem=256gb
#SBATCH --time=14-00:00:00
#SBATCH --output=sparc_service_%j.log

module purge
module load apptainer

RIVA_SIF="/blue/your_group/sparc_project/containers/riva_server.sif"
BRIDGE_SIF="/blue/your_group/sparc_project/containers/websocket_bridge.sif"
MAS_SIF="/blue/your_group/sparc_project/containers/mas_server.sif"

# Launch Services in Background
echo "Starting Riva..."
apptainer exec --nv ${RIVA_SIF} riva_start.sh &
```

```
sleep 20

echo "Starting Bridge..."
apptainer exec ${BRIDGE_SIF} riva-websocket-gateway --riva-uri=localhost:50051 --
port=8080 &

echo "Starting MAS..."
apptainer exec --nv ${MAS_SIF} uvicorn main:app --host 0.0.0.0 --port 8000 &

wait
"""
with open("sparc_production.slurm", "w") as f:
    f.write(script_content.strip())
print("Generated sparc_production.slurm")

generate_production_script()
```

---

## Summary

This notebook provides a complete containerization and deployment workflow for the SPARC backend:

1. **Multi-Stage Docker Build:** Minimizes container size and security vulnerabilities by separating build and runtime stages.
2. **Local Development with Podman:** Uses Podman Pods to create a shared network namespace that simulates the production environment on a developer's machine, allowing all services to communicate via localhost.
3. **Production Deployment on HiPerGator:**
  - Converts Docker images to Apptainer SIF format for HPC compatibility
  - Uses SLURM job scheduling to allocate GPU resources
  - Launches services as persistent background processes
  - Includes Riva Speech Server, WebSocket-to-gRPC bridge, and Multi-Agent System
4. **Key Components:**
  - **Riva Server:** Handles speech processing tasks
  - **WebSocket Bridge:** Enables communication between Riva and Unity frontend via WebSocket
  - **MAS Server:** Runs the trained agents for handling interactions

The deployment is fully containerized, reproducible, and scalable on HPC infrastructure.